

CS 5/7330

NoSQL : Document Databases / MongoDB

Motivation: Semi-structured data

- Relational model is very rigid in terms of structure
- For example, consider this table:

SSN	Name	Phone-#
123456789	John Doe	123-4567
234567890	Mary Doe	456-8901

- But what if someone have more/less than one phone?

Motivation: Semi-structured data

- No phone

SSN	Name	Phone-#
123456789	Brad Doe	NULL
234567890	Mary Doe	456-8901

- What can NULL means?
 - Brad has no phone
 - Brad has a phone but I don't know his number
 - Possible ambiguity

Motivation: Semi-structured data

- Many phones

SSN	Name	Phone-#
123456789	Tai-Kwan Doe	555-5555
123456789	Tai-Kwan Doe	689-0777

- Cause duplication
- Anomaly and/or inefficiency



Motivation: Semi-structured Data

- Many problems can be resolved by normalization (breaking into multiple tables)
- But queries will require potentially many joins
 - Can be inefficient
- Solution:
 - Semi-structured data

Semi-structured Data

- Some data have a bit of a structure
- E.g. item have a set of attributes
- However they can still differs
 - some attributes may be present in some items but not all of them
 - Some attributes for an item may have multiple values
 - ... or even multiple possible types
- So new models are proposed. Examples:
 - XML
 - JSON

JSON

- JavaScript Object Notation
- Originally decided as a way for information interchange
 - Text-based
 - Self describing
 - The “schema” is described within the object
- While the schema for an object is flexible, the way of describing the schema is fixed
 - Thus one can write standard tool(s) to interpret any JSON object, without knowing its schema beforehand

JSON

- Each JSON object is represented as a set of name/value pairs
 - {"name":"John", "age":30, "car":null}
- Name is a string in quotes (“name”, “age”, “car”)
- Each name has a value associated with it.
- Valid value types:
 - String
 - Number
 - Array
 - Boolean
 - Null
 - Another JSON object

JSON (Example of complex case)

<https://www.sitepoint.com/twitter-json-example/>

General idea of document databases

- Database store objects
 - Objects are known as documents
 - Name/value pairs are usually termed as key/value pairs
- Objects have unique IDs
 - Typically as Entity in E-R Model
- Relations between objects can be enforced by
 - Nested Objects/Documents (embedding)
 - Key/value pairs where values are IDs of documents to be related (references)

General idea of document databases

- Other properties
 - Data are often grouped into buckets (like tables)
 - Provide indexes for document attributes
 - Provide query language support to retrieve (part of) a document

Modeling using Document Databases

- One-to-one relationship
- Either embedding or relationship will work
- Example : (assume patron and address is a 1-1 relationship)

```
// patron document
{
  _id: "joe",
  name: "Joe Bookreader"
}

// address document
{
  patron_id: "joe", // reference to patron document
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}
```

```
{
  _id: "joe",
  name: "Joe Bookreader",
  address: {
    street: "123 Fake Street",
    city: "Faketon",
    state: "MA",
    zip: "12345"
  }
}
```

Modeling using Document Databases

- One-to-many relationship
- Example : (assume patron and address is a 1-m relationship)
- Embedding

```
{
  "_id": "joe",
  "name": "Joe Bookreader",
  "addresses": [
    {
      "street": "123 Fake Street",
      "city": "Faketon",
      "state": "MA",
      "zip": "12345"
    },
    {
      "street": "1 Some Other Street",
      "city": "Boston",
      "state": "MA",
      "zip": "12345"
    }
  ]
}
```

Modeling using Document Databases

- One-to-many relationship
- Example : (assume publisher and books is a 1-m relationship)
- relationship

```
{
  name: "O'Reilly Media",
  founded: 1980,
  location: "CA",
  books: [123456789, 234567890, ...]
}

{
  _id: 123456789,
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English"
}

{
  _id: 234567890,
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English"
}
```

Modeling using Document Databases

- What you should AVOID
- Duplication of data
- Same problem as in relational databases

```
{  
  title: "MongoDB: The Definitive Guide",  
  author: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  publisher: {  
    name: "O'Reilly Media",  
    founded: 1980,  
    location: "CA"  
  }  
}  
  
{  
  title: "50 Tips and Tricks for MongoDB Developer",  
  author: "Kristina Chodorow",  
  published_date: ISODate("2011-05-06"),  
  pages: 68,  
  language: "English",  
  publisher: {  
    name: "O'Reilly Media",  
    founded: 1980,  
    location: "CA"  
  }  
}
```

Modeling using Document Databases

- One thing you should never do
- Whenever there is a relationship between two objects, using relationship should only be from one object to the other, BUT NOT BOTH WAYS
 - Why not?
- Which way to go is database and application dependent

```
{ _id = "_1";  
  name = "john doe"  
  address = "_add1"  
}  
{  
  _id = "_add1"  
  address = "SMU Blvd"  
  patron_id = "1"  
}
```


Modeling using Document Databases

- Document models tends to works well with binary relationship.
- For relationship with multiple entity (e.g: supplier-restaurant-food), it may need either:
 - Deep embedding
 - Creating separate objects to store the relationships (in terms of E-R modelling, treating relationship as entity)

Example: MongoDB

- Document-based database system
- Basic unit of storage: Document
 - Very similar to JSON
 - With some additional data type support (e.g. date, binary data, code)
 - Each document needs to have an “_id” field (can be auto-generated)
 - Each document has a size limit : 16 MB (in Mongo DB 4.0)
- Collections: A group of Documents
 - No requirement that each document have the same schema
 - But there are advantages (in terms of implementation, and potential efficiency)
 - One can name sub-collections of a collection
- Database: A group of collections
 - Typically an application will use one database

MongoDB: insert

- Insertions:
 - `db.collection.insertOne({...})`
 - `db.collection.insertMany([{...}, {...}, ..., {...}])`
- Notice that the only automatic check is unique `_id`
 - NO notion of referential integrity
 - NO checking of name of keys (since we do not assume any scheme [even within a collection])
 - It's the (virtually all) developer's responsibility to ensure everything works.

MongoDB: Query -- selection

- Basic method: `find()`
- `db.collection.find({<selection>}. {<projection>})`
 - Describe the documents to be retrieved
 - Each of this has the same format as a document (but without the `_id`).
- `db.collection.find({})`
 - Return all documents in the collection
- `db.collection.find({ "name" : "john doe" })`
 - Return all documents that have key value pair (name : john doe)
- `db.collection.find({ "name" : "john doe", "age" : 30 })`
 - Return all documents that have BOTH key value pairs
 - Logical AND
- Note:
 - Value in the key/value pair in find must be a constant.
 - Need other constructs to relate different documents from the same/different collections

MongoDB: Query -- projection

- `db.collection.find({"name" : "john doe"}, {"name" : 1, "age": 1})`
 - Return all documents that have key value pair (name : john doe), and output the name and the age
- `db.collection.find({"name" : "john doe"}, {"name" : 0, "age": 1})`
 - Return all documents that have key value pair (name : john doe), and output age, but not the name
- Note:
 - `_id` : is output by default
 - Can suppress it by including `{"_id" : 0}` in the projection clause

MongoDB: Query – selection clauses

- `db.collection.find({“name” : “john doe”, “age” : {“$gte”: 18}})`
 - Return all documents that have name john doe, and is at least 18 years old (gte = greater than or equal to)
 - Clause: a document, the key is a reserved word starting with “\$”, to denote special condition or functions,
- `db.collection.find({“name” : “john doe”, “age” : {“$gte”: 18, “$lte” : 22 } })`
 - Return all documents that have name john doe, and is between 18 and 22 (lte = less than or equal to)

MongoDB: Query – OR / NOT

- `db.collection.find({“dept” : “CS”, {$or, [{“age” : {“$gte”: 18}}, {“gpa” : 4.0}] } })`
 - Return all documents that have dept = “CS” and either age is ≥ 18 , or gpa = 4.0
 - Notice that the term following or is an list (array) of conditions.
- `db.collection.find({“dept” : “CS”, {$not, {“gpa” : 4.0} } })`
 - Return all documents that have dept = “CS” and either age is ≥ 18 , or gpa $\neq 4.0$

MongoDB: Query – NULL values

- BE VERY CAREFUL!
- Suppose your collection has the following documents:
 - { “_id” : “1”, “name” : “john doe”, “phone” : null }
 - { “_id” : “2”, “name” : “jack doe”, “age” : 28 }
- `db.collection.find({“phone” : null})` will return BOTH documents
 - Null matches either “having the value null”, and “no such key exists in the document”
 - Use the “\$exists” clause to compensate

MongoDB: Query – Arrays

- { “_id” : “1”, “name” : “john doe”, “x” : [1, 3, 5, 11]}
- The following queries will return this document
 - db.collection.find({“x” : 5})
 - db.collection.find({“x” : {“\$gte”, 4, “\$lte”, 11}})
 - db.collection.find({“x” : {“\$gte”, 7, “\$lte”, 9}})
 - Why?
Notice that matching an array means that every clause has to be matched, but each clause can be matched by a **different** element in the array
 - Use \$elemMatch to overcome this

MongoDB: Query – Embedded documents

```
{ "_id" : "1", "name" : "john doe",  
  "course" : { "cid" : "CS 5330", "grade": "A" } }
```

- `db.collection.find({"course" : {"cid" : "CS 5330", "grade": "A"} })`
 - will return the document
- `db.collection.find({"course" : {"cid" : "CS 5330"} })`
 - Will NOT
 - For embedding documents, require perfect match ("`_id`" can be left out)
- `db.collection.find({"course.cid" : "CS 5330"} })`
 - Will return the document again
 - Use dot notation

MongoDB: Query – More complicated issues

(See MongoDB manual / books for details)

- Matching array of embedded documents can be very tricky
- No inherent way to “join” multiple collections using the query language
 - Need to write program and using API
 - There is an option of a “\$WHERE” clause where you can embed code inside a query statement

Mongo DB : Updates and Delete

- `db.collection.updateOne()`, `updateAll()`, `replaceOne()`, `replaceAll()`
 - Two parameters
 - First one is the query condition to specify which document(s) in the collection is to be updated.
 - Second one specify how the document is to be updated / or what to replace the document with
 - For replacement, the “_id” will not change, but everything else will be replaced.
 - Also an “upsert” option: try to update, if the document to be updated is not found, then insert a new document.

Other aspects

- Indexing – per collection basis, similar to relational database indices
 - Specific type of indices for specific data types (e.g. text, geospital)
- Aggregation – group by, and computation based on groups
- Replication – allow for duplication of data