

# CS 5/7330

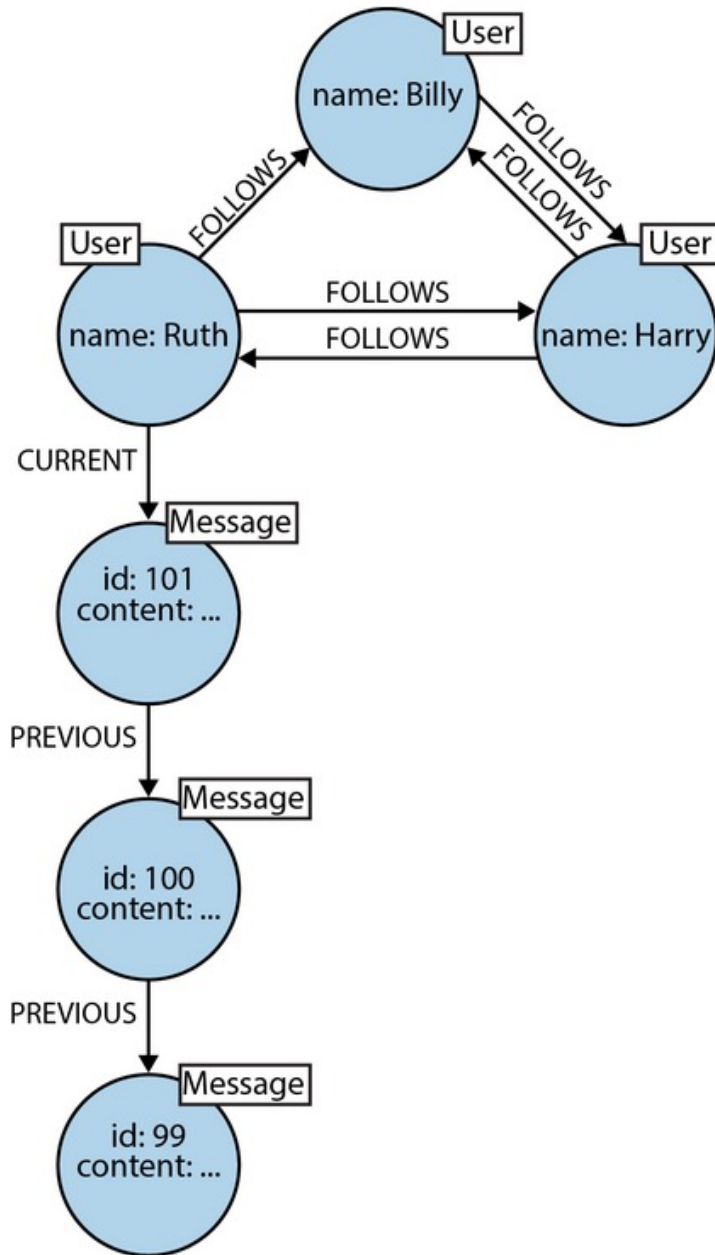
Graph Databases

# Graphs

- Graphs
  - Vertices and edges
  - Vertices (may) have (multiple) labels/attributes
  - An edge connects a pair of vertices
    - Can be directed (or not)
    - Can have (multiple) labels/attributes
    - May have multiple edges between any pairs of vertices (not common, but can be useful)
- Graphs can be used to model:
  - Relationships
  - Maps
  - Networks

# Graphs

Example:



*From: Eifrem, Robinson,  
Webber, "Graph  
Databases", 2<sup>nd</sup> edition*

Figure 1-2. Publishing messages

# Graphs

Example:

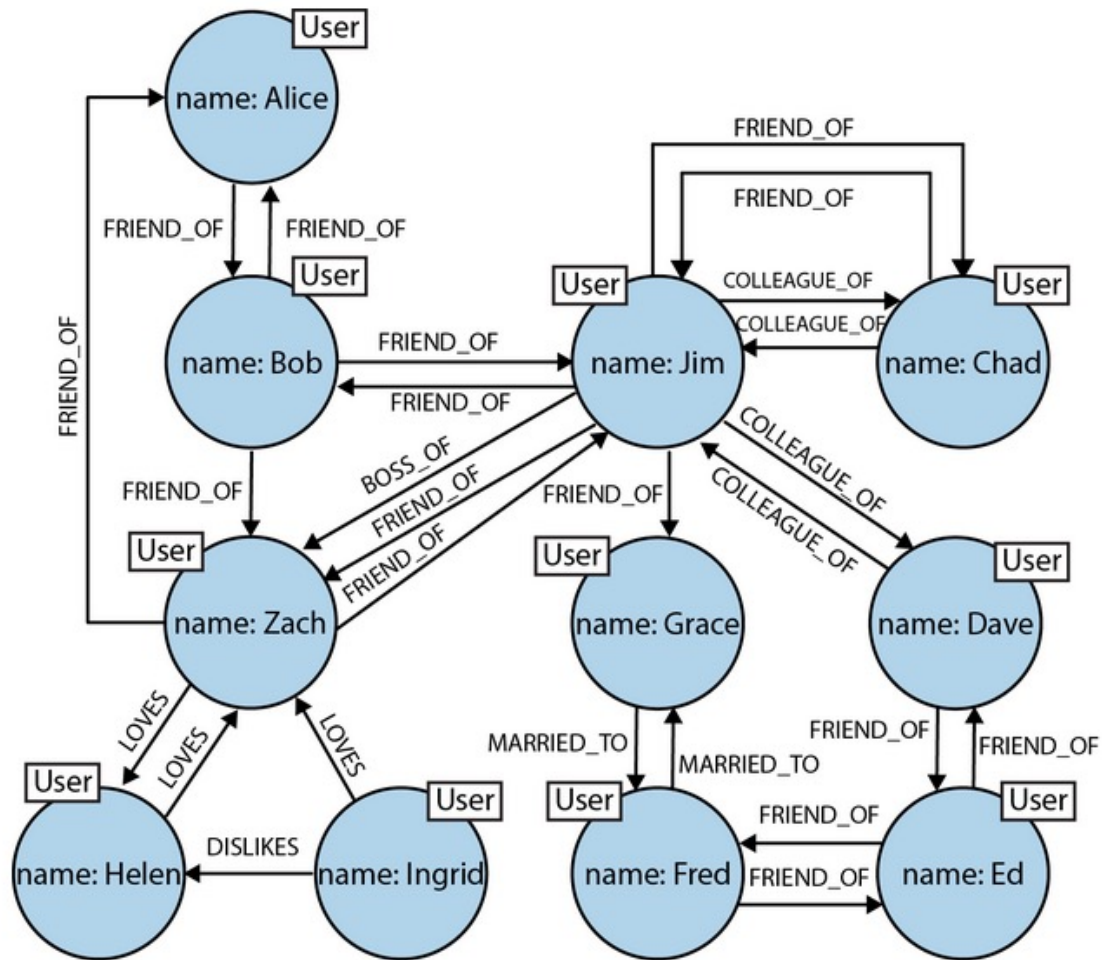


Figure 2-5. Easily modeling friends, colleagues, workers, and (unrequited) lovers in a graph

*From: Eifrem, Robinson, Webber, "Graph Databases", 2<sup>nd</sup> edition*

# Different types of graphs

Table 1. Various Graph Structures

Graph structure	Description
<b>Undirected/directed graphs</b>	All relationships in an undirected graph are symmetric. On the other hand, in a directed graph, each edge $e$ from vertex $src$ to vertex $des$ is a directed tuple, such that (1) edge $e$ is an <i>in-edge</i> of $des$ and an <i>out-edge</i> of $src$ , (2) vertex $des$ is the $src$ 's <i>out-neighbor</i> and vertex $src$ is the <i>in-neighbor</i> of $des$ , and (3) the number of <i>incoming/outgoing</i> edges of a vertex determines its ( <i>in/out</i> ) degree.
<b>Labeled graphs</b>	Vertices and edges are tagged with scalar values (labels or types) that may represent either their roles in different application domains or some attached metadata.
<b>Attributed graphs</b>	A variable list of attributes as ( <i>key, value</i> ) pairs are attached to vertices and edges, representing their properties. It is suitable for social networking sites that involve social interaction of individuals.
<b>Multigraphs</b>	Multiple edges (even with the same labels) between the same two vertices as well as self-loops are allowed.
<b>Hypergraphs</b>	These graphs can represent N-ary relationships through hyperedges that can connect any number of vertices (Bretto 2013). An undirected hyperedge can be represented by a subset of nodes (vertices) (Berge 1973); and a directed hyperedge can be represented by a tuple (ordered set) of nodes (Boley 1992) or head-and-tail sets (Gallo et al. 1993; Nguyen and Pallottino 1989). Hypergraphs are used to represent complex relationships in areas such as Artificial Intelligence. HyperGraphDB (Iordanov 2010) is based on hypergraphs and supports N-ary relationships by representing a link as a tuple (ordered set) of nodes. Figure 4 shows a simple hypergraph.
<b>Nested graphs</b>	Each vertex, in turn, can be a graph. At this time, no store supports nested graphs.

# Typical Operations for graph databases

- On-line graph navigations:
  - Explore a (typically small) subset of vertices
  - Require fast response time
- Two types
  - Path query
    - Whether a path exists between two given vertices
    - Conditions on the edges along the path
    - Can also be given only one vertex and find all the other vertices that can be reached/can reach it from the query vertex

# Typical Operations for graph databases

- On-line graph navigations:
  - Explore a (typically small) subset of vertices
  - Require fast response time
- Two types
  - Pattern matching query
    - Find all subgraphs that is isomorphic to the query subgraph
    - General graph isomorphism is NP-complete
    - Tree isomorphism has polynomial time solutions

# Typical Operations for graph databases

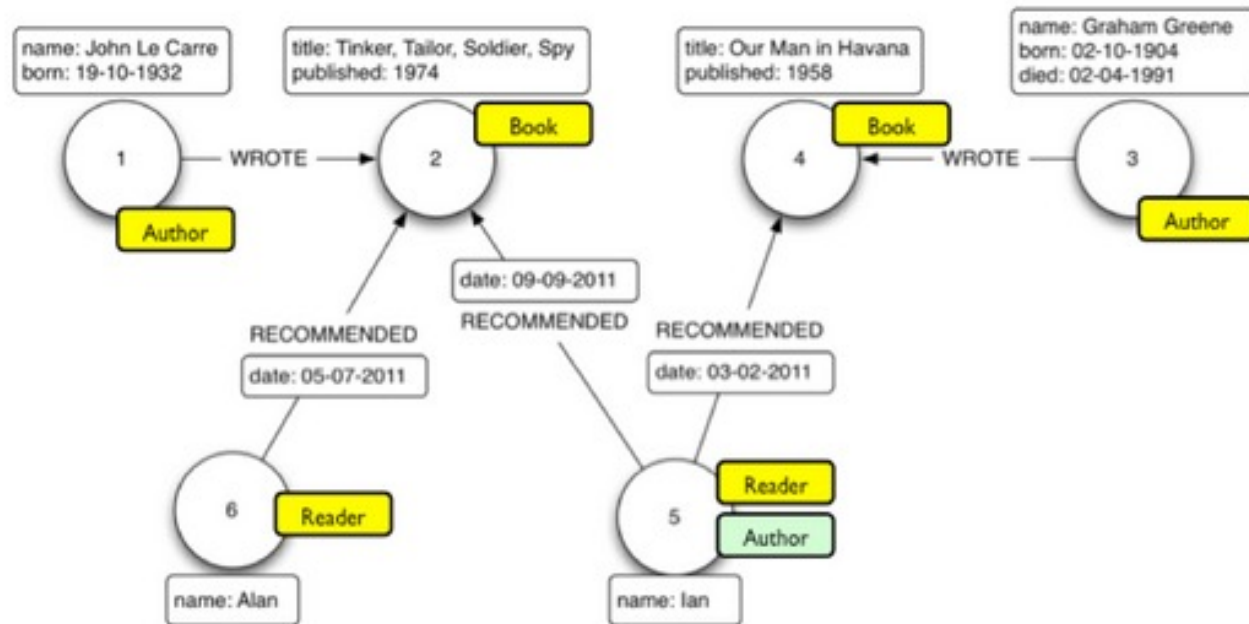
- Off-line analytical graph computation
  - Analyze large (number of) graphs
  - Typically iterative computation (not map-reduce)
  - Require highly tuned graph packages
  - Aim for high throughput, not low latency



# Example : Neo4j

- Based on the Labelled Property Graph Model
- 4 major component of a graph
  - Nodes – generally for entity
  - Relationships (edges) – generally for relationship between entities
  - Properties – key/value pairs that is associated with nodes and/or relationships
    - Like document model, no need for fixed schema
  - Labels: Assigned to nodes (notice that, just labels, not key/value pairs)

# Example : Neo4j



The labeled property graph model

From, Jerome Baton, “Learning Neo4j 3.x”, 2<sup>nd</sup> edition

# Neo4j Data Model

- Work best for binary relationship between entities
- May need extra node for n-ary relationships
- Challenge of whether making something a property or a node
  - Make graphs simpler vs. query speed
  - Depends on whether it is faster to traverse a graph vs. search node/edge based on properties

# Neo4j Data Model

- General mapping from E-R model to Graph Model
  - Entity -> Nodes (with labels serve as names)
  - (Binary) Relationship -> Edges
  - Attributes -> Properties

# Cypher

- Common database language for graph databases
- Based on patterns and pattern matching
- Basic ideas:
  - Database : a set of nodes and edges
  - Commands to create those nodes and edges
  - Queries: specifying a pattern of edge/node/paths/subgraphs and try to match it.

# Cypher

Creating a node:

- `CREATE (u: User {username: "greg", age: "33"})`
  - Return a User (label) node (u) that has a property (name/"greg", age/33)
  - The variable u is used as a reference to the node
  - You can create multiple nodes (separate each node with a comma) in a statement
- `CREATE (u: User:Admin {username: "greg", age: "33"})`
  - A node can have multiple labels
  - This node has two labels, User, Admin.
    - No notion of sub-labels.

# Cypher

## Finding a node

- MATCH (a:User), (b:System)  
WHERE a.username="greg" AND b.sysname = "UNIX"  
RETURN a, b
- Notice that if multiple nodes matches they will all be returned (like a Cartesian Product)

## Projections

- MATCH (a:User {name: "John Doe"})  
RETURN a.age, a.gender

# Cypher

## Patterns used for queries

- (a)
  - matching any node.
  - a is the variable name that use to refer to the node in the rest of the query
- (a) - - > (b)
  - Match any pair of nodes that go from a to b
- (a) < - - (b)
  - Match any pair of nodes that go from b to a
- (a) - - > (b) - - > (c)
  - Match three nodes that form a path from a to b to c



# Cypher

## Patterns used for queries

- (a: User)
  - matching any node with the label user.
  - a is the variable name that use to refer to the node in the rest of the query
- (a: User:Admin)
  - Matching nodes with BOTH labels
- (a:User|Admin)
  - Matching nodes with either label

# Cypher

## Patterns used for queries

- (a: {name:"John", age:25})
  - matching any node with the following properties.
- (a) -> [{length : 30}]-> (b)
  - Matching a path where the edge has the property specified

# Cypher

## Building an edge

- MATCH (a:User), (b:System)  
WHERE a.username="greg" AND b.sysname =  
"UNIX"  
  
CREATE (a)-[pr:LOGIN]->(b)  
SET pr.date ="11/11/2020"  
  
RETURN pr
- Create an edge from node a to node be, with the label  
"LOGIN" and key/value pair (date/"11/11/2020")

# Cypher

## Constraints:

- Uniqueness Constraint:
  - `CREATE CONSTRAINT ON (a:USER) ASSERT a.username IS UNIQUE`
  - Ensure uniqueness of username for all USER nodes
- Property Existence Constraint:
  - `CREATE CONSTRAINT ON (a:User) ASSERT exists(a.username)`
  - Ensure every node of User have username property