

CS 5/7330

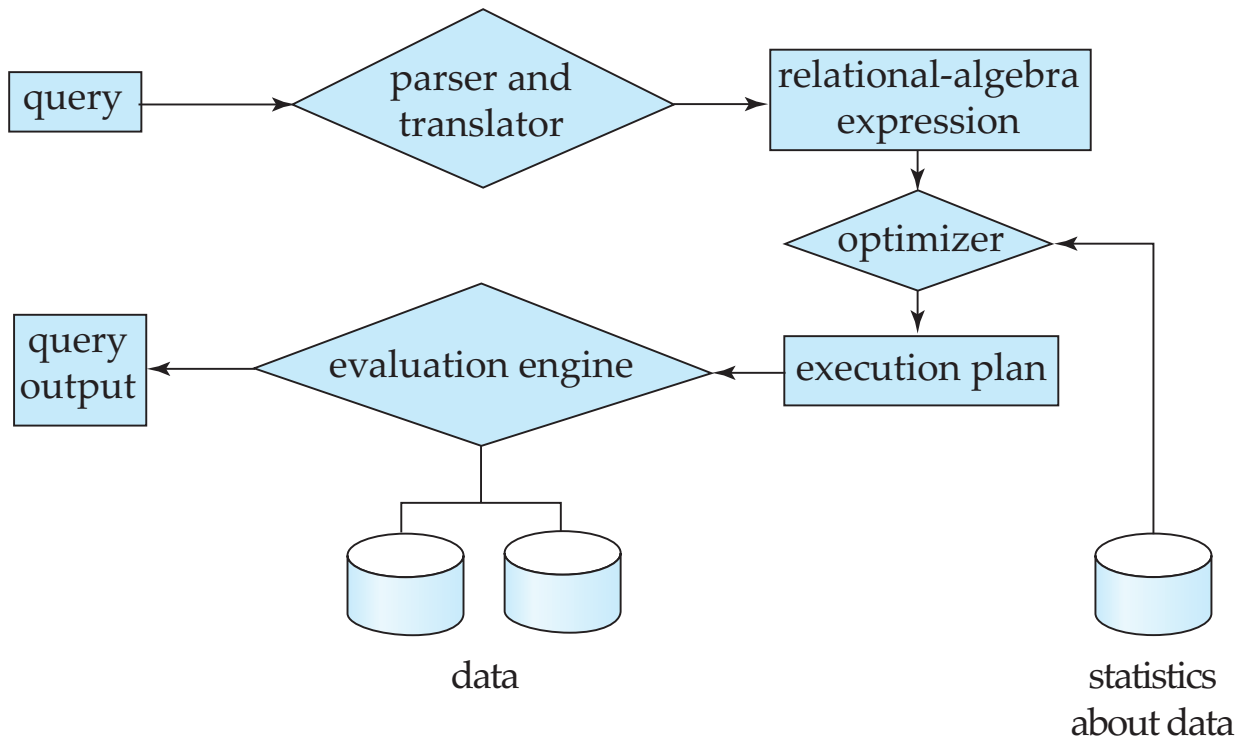
Query Processing / Optimization

Query Processing / Optimization

- Study how database process query internally
 - i.e. Once a database system received an SQL query, what happens (until the database return the results)
- Why study?
 - Database project manager – understand whether the queries being written will likely be executed effectively by the database
 - Database administrator – able to restructure database / provide hints to the database system to speed up queries
 - Database developer – you may be hired by a database systems company to build the next version of a DBMS

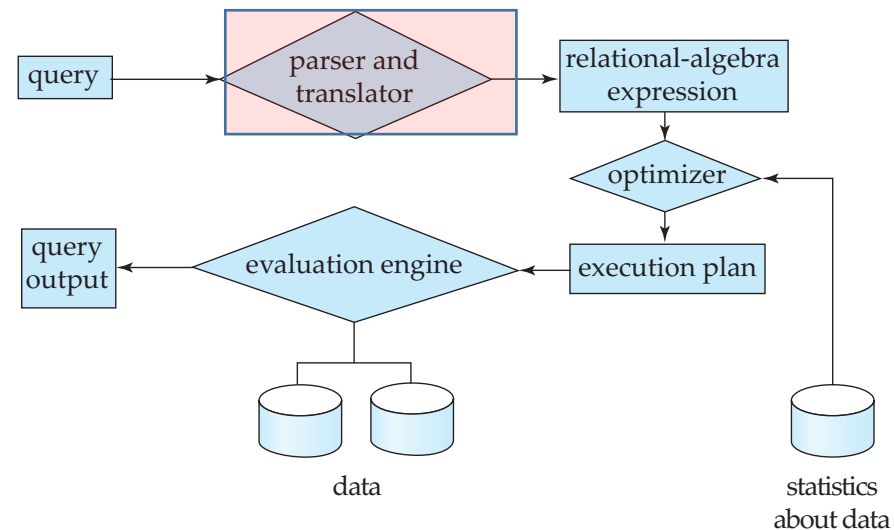
Query Processing

- What happen when a DBMS received a query



Query Processing

- Parser and translator
- Read the query
- Check for syntax
- Collect all relevant information about tables involved
- Break the query down into a set of basic operations (relational algebra + others)



Query Processing

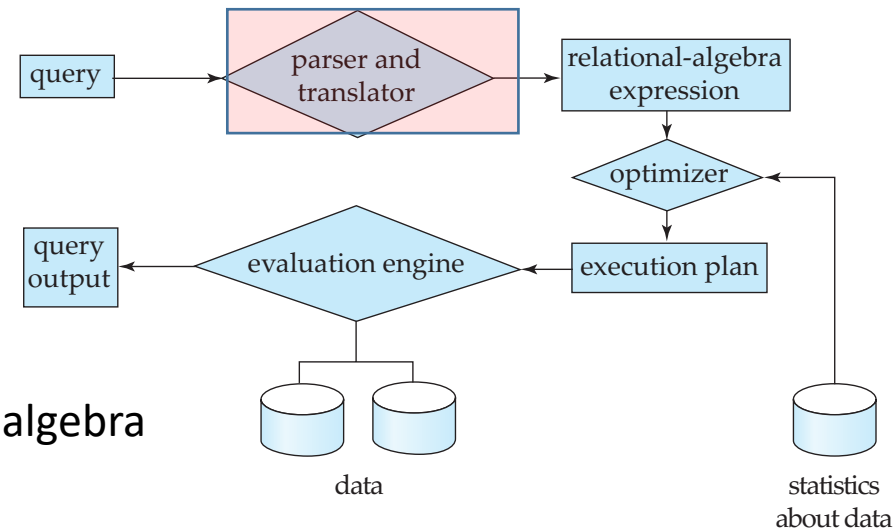
- Recall

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

Projection in relational algebra

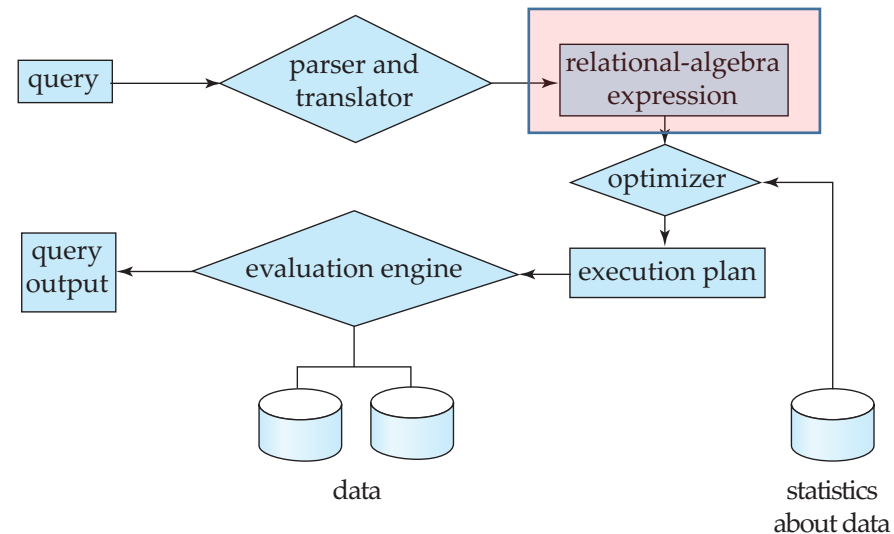
Cartesian Product in relational algebra

Selection in relational algebra
(remember, join = Cartesian
product + selection)



Query Processing

- Query is converted to a list of operators (σ , Π , \bowtie , Group, Order)
- Thus running the query becomes execution of a list of operators




Query Processing

- Example

SELECT ssn, age
FROM Instructor
WHERE age > 25

Becomes

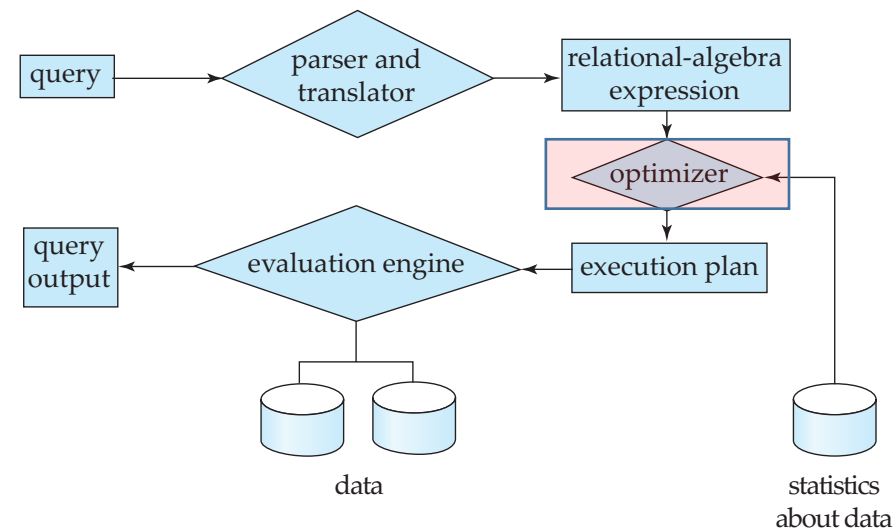
1. $A = \sigma_{\text{age} > 25} (\text{Instructor})$
2. $\text{Result} = \prod_{\text{ssn, age}} (A)$



Query Processing

Optimizer

- For each operation, determine how it will be executed
- Determine the order of operations
- Other tasks (to be discussed later)



Query Processing

Execution plan

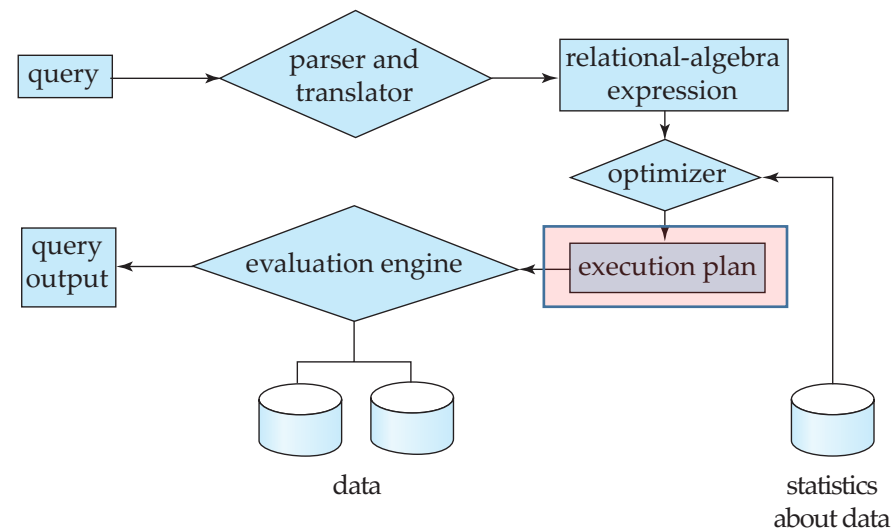
- Example:

1. $A = \sigma_{\text{age} > 25}(\text{Instructor})$

- Use the secondary index on A to retrieve the tuples
- Do not store A on to the disk

2. $\text{Result} = \Pi_{\text{ssn}, \text{age}}(A)$

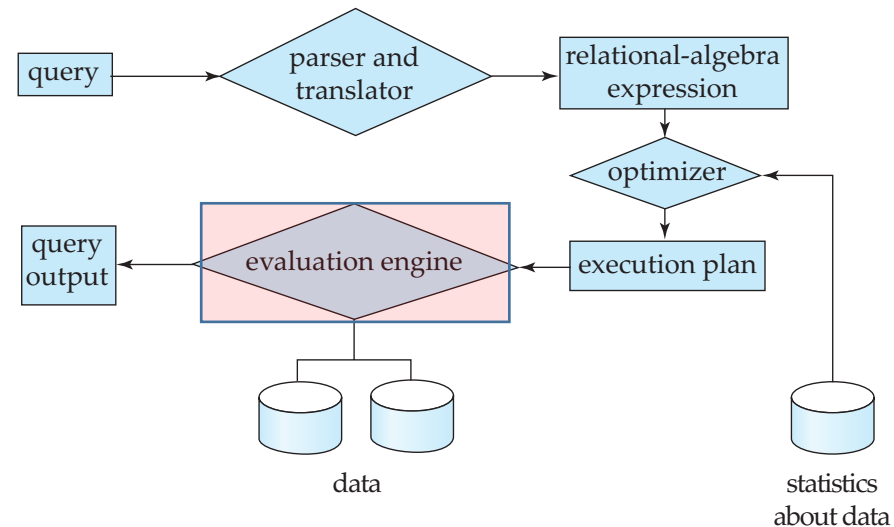
- Directly pipe the result from 1 to execute the project
- Just pick the attributes and output them



Query Processing

Evaluation engine


- Actual execute the plan
- Return the tuples to the output
 - Can be on screen, as a file, or as a stream through a network



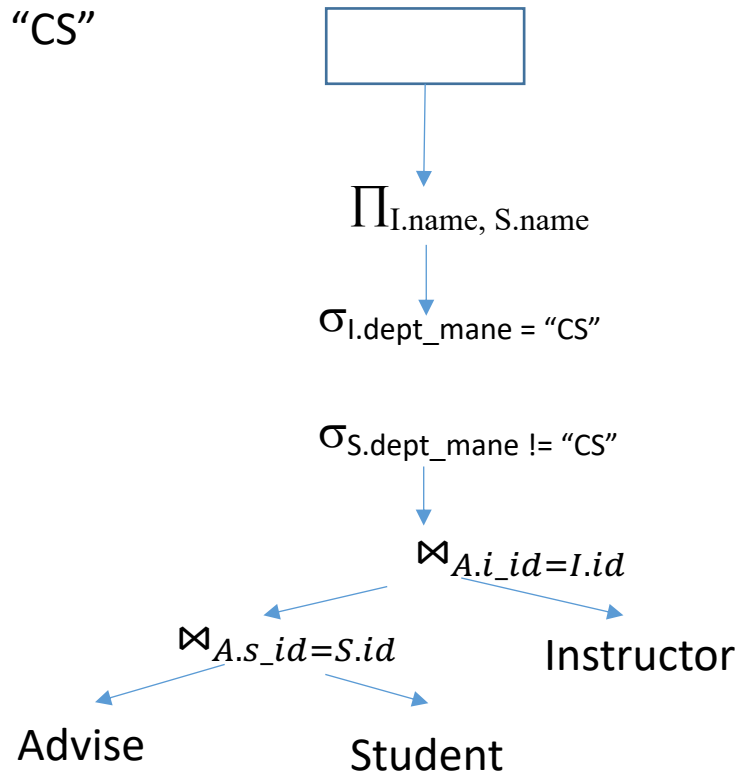
Query Processing

- Notice that more often, a query can be represented as a parse tree:

```
SELECT ssn, age
FROM Instructor
WHERE age > 25
```



```
SELECT I.name, S.name
FROM Instructor I, Advise A, Student S
WHERE A.i_id = I.ID AND S.id = A.s_id
AND I.dept_name = "CS" AND
S.dept_name != "CS"
```



Query Processing

- Query Processing/Optimization in terms of parse trees:
 - Parsing: Generate parse tree for the query
 - Notice that a query may generate MANY (TOO MANY) valid parse trees
 - Each node correspond to an operation
 - Evaluation of query is from bottom to top (think post order)
 - Optimization:
 - Decorating the nodes of a parse
 - For each node, determine how that operation is to be executed
 - Select the best parse tree
 - For all the parse tree generated, pick the one that will execute the query fastest

Query Processing / Optimization

- Challenges
 - Combinatorics explosion
 - There are many valid parse tree for a query
 - There are many possible ways to decorate each node
 - The total number of options grows exponentially (since most combinations of parse tree and decorations are valid)
 - Limited time
 - You do not want to spend 2 hours optimizing a query such that you save 2 seconds.
 - Information needed
 - A lot of information is needed for the optimizer to work properly
 - Recall the fundamental problem of query optimization?

Measuring cost of a query

- Many factors contribute to time cost
 - *disk access, CPU, and network communication*
- Cost can be measured based on
 - **response time**, i.e. total elapsed time for answering query, or
 - total **resource consumption**
- We use total resource consumption as cost metric
 - Response time harder to estimate, and minimizing resource consumption is a good idea in a shared database
- We ignore CPU costs for simplicity
 - Real systems do take CPU cost into account
 - Network costs must be considered for parallel systems
- We describe how estimate the cost of each operation
 - We do not include cost to writing output to disk

Measuring cost of a query

- Disk cost can be estimated as:
 - Number of seeks/rotates * average-seek-cost
 - Number of blocks read * average-block-read-cost
 - Number of blocks written * average-block-write-cost
- For simplicity we just use the **number of block transfers** *from disk* and the **number of seeks** as the cost measures
 - t_T – time to transfer one block
 - Assuming for simplicity that write cost is same as read cost
 - t_S – time to move from one block to a non consecutive block
 - If on the same track, then rotate
 - If on different track, then seek + rotate
 - Cost for b block transfers plus S seeks
$$b * t_T + S * t_S$$
- t_S and t_T depend on where data is stored; with 4 KB blocks:
 - High end magnetic disk: $t_S = 4$ msec and $t_T = 0.1$ msec
 - SSD: $t_S = 20$ -90 microsec and $t_T = 2$ -10 microsec for 4KB

Measuring cost of a query

- Required data may be buffer resident already, avoiding disk I/O
 - But hard to take into account for cost estimation
- Several algorithms can reduce disk IO by using extra buffer space
 - Amount of real memory available to buffer depends on other concurrent queries and OS processes, known only during execution
- Worst case estimates assume that no data is initially in buffer and only the minimum amount of memory needed for the operation is available
 - But more optimistic estimates are used in practice

Query processing for individual operations

- Selection
- Joins
- Projection
- Ordering
- Group By

Selection

- $\sigma_{\text{condition}}$ (Table)
- A few considerations
 - Condition
 - Can be “attribute = value” (equality)
 - Or “attribute <= value” (or <, >, >=) (range/comparison)
 - Single condition or multiple condition
 - E.g. (x = 1 AND y = 3), (x = 3 OR y = 4)
 - Attribute
 - Primary key or not
 - Organization
 - Index available?
 - What kind of index?

Selection

- Basic case: File scan / sequential scan
- Just read the whole file block-by-block from beginning to end and check if each tuple satisfy conditions
 - Cost estimate = $b_r * t_T + b_s * t_S$
 - b_r denotes number of blocks containing records from relation r
 - B_s denotes number of tracks that store the table
 - If selection is on a key attribute, can stop on finding record
 - Average cost = $(b_r/2) * t_T + ? * t_S$
 - ? Is harder to predict, is roughly $\max(1, b_s/2)$
 - Linear search can be applied regardless of
 - selection condition or
 - ordering of records in the file, or
 - availability of indices
- Note: binary search generally does not make even if data is sorted except when there is an index available, as each step will require a rotation

Selection

- If you have an index: Index scan
 - Use an index to search.
 - Attribute of index needs to match condition
 - Also hash table is not useful for range queries
- Assume that the query return t tuples, stored in b_r blocks (notice that $t > b_r$ – often by at least one order of magnitude)
 - $t = 1$ if equality search on an unique attribute
- Clustering index:
 - Cost = time for searching the index + $b_r * t_T + \alpha * t_S$
 - α is the number of tracks that contains retrieved tuples (if $t = 1$, $\alpha = 1$)
 - α grow slowly (if at all) with t (since data is clustered)

Selection

- If you have an index: Index scan
 - Use an index to search.
 - Attribute of index needs to match condition
- Assume that the query return t tuples, stored in b_r blocks (notice that $t > b_r$ – often by at least one order of magnitude)
 - $t = 1$ if equality search on an unique attribute
- Non-Clustering index:
 - Worst case Cost = time for searching the index + $t * (t_I + t_S)$
 - It may be faster if the DBMS read all the index record and determine which tracks the data are stored before fetching the record
 - Even this can be infeasible
 - For example, if there is an ORDER BY clause and there are too many records being selected to fit in main memory
 - Cannot bring them all in to sort

Selection

- Conjunction: $\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \wedge \theta_n}(r)$
 - Now multiple indices may be available
 - Option 1: Single index
 - Use one of the available index
 - Read in tuples into main memory, then apply other conditions
 - Some cost estimation as before, only with t (and br) are tuples/blocks that contain tuples that satisfy the condition that the index is correspond to
 - Option 2: Clustering multiple-attribute index (if available)
 - Option 3: Intersection of identifiers
 - Consider all secondary indices that associate with an attribute in the condition
 - Query the index, read in all index records into main memory (do not go to database yet)
 - Only select pages that are in the answers for all indices
 - Read the tuples and subsequently apply other conditions

Selection

- Disjunction: $\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)$.
 - Now multiple indices may be available
 - Option 1: Intersection of identifiers
 - Assume indices present for ALL attributes in the condition
 - Query the indices, read in all index records into main memory (do not go to database yet)
 - Only select pages that are in at least one of the indices
 - Usually is expensive
 - Notice that it doesn't work if indices for some attributes in the condition is not available
 - Option 2: Sequential scan, apply condition when tuple is read

Joins

- $R \bowtie_{\text{cond}} S$
- Most common case: condition is an equality condition between attributes of R and S
 - We call this equi-join
 - Example: linking an attribute with its foreign key
- With equi-join there are a lot of options
- With non equi-join there are very few

Joins – Nested loop

- The naïve algorithm
- $R \bowtie_{\text{cond}} S$
- **for each tuple t_r in R do begin**
 - for each tuple t_s in S do begin**
 - test pair (t_r, t_s) to see if they satisfy the join condition *cond*
 - if they do, add $t_r \bullet t_s$ to the result.
 - end**
- end**
- R is called the outer relation (outer loop), S is the inner relation (inner loop)
 - Either relation can be in the outer loop
 - Flipping R and S will give you the same results
- Works for any condition

Joins – Blocked Nested loop

- Modify the algorithm for secondary storage
- **for each block b_r in R do begin**
 - Read b_r from disk into main memory
 - for each block b_s in S do begin**
 - Read b_s from disk into main memory
 - do a nested loop for each pair (t_r, t_s) [$t_r \in b_r, t_s \in b_s$] to see if they satisfy the join condition *cond*
 - if they do, add $t_r \bullet t_s$ to the result.
 - end**
- end**

Joins – Blocked Nested loop

- Running time depends on amount of main memory buffers available
 - Need at least 2. 1 for R and 1 for S
- Case 1: Buffer large enough to hold all blocks for both tables
 - Read both tables into main memory and then loop all the tuples inside
 - Cost = cost for sequential scan for R + cost for sequential scan for S

Joins – Blocked Nested loop

- Running time depends on amount of main memory buffers available
 - Need at least 2. 1 for R and 1 for S
- Case 2: Minimum number of buffers (ctd)
 - Separate cost into page access and seeks
 - For page access, each block in the outer loop need to be read once
 - For inner loop, each block has to be read once for each block of the outer loop
 - Number of block access = $b_r + b_r * b_s$
 - Question: given R and S, which table should be in the outer loop

Joins – Blocked Nested loop

- Running time depends on amount of main memory buffers available
 - Need at least 2. 1 for R and 1 for S
- Case 2: Minimum number of buffers
 - For seeks/rotate
 - Each page in the outerloop need to be seeked
 - Assume query is not being interrupted, the inner loop is being read in consecutively
 - So only the minimum number of seeks
 - Number of seeks = $b_r * (\text{number of seeks to read S})$
 - However, it is more than likely that there will be interruptions
 - For example: the disk head may have moved pass while the joining is in process
 - Worst case scenario: $b_r * b_s$
 - Which table should be in the outerloop?

Joins – Blocked Nested loop

- Case 3: Enough buffer to fit either R and S (plus k buffers for the other table), but not both
 - (Assume the buffer fits table S)
 - Read S into main memory
 - Then read R into main memory (in steps, because not enough memory to read it all at once)
 - Join pairs of tuples in R and S
 - Cost = cost of reading S (sequential scan) + cost of reading R
 - Depends on whether the query get interrupted, it can be as little as the same as sequential scan or as much as $\text{ceiling}(b_r / k)$, where k is the number of buffers allocated to R

Joins – Blocked Nested loop

- Running time depends on amount of main memory buffers available
 - Need at least 2. 1 for R and 1 for S
- Case 4: None of the above (k buffers available between the two tables)
 - We will need to assign buffers to each table
 - Assume we assume k' buffers to R
 - Then $k-k'$ buffers to S
 - **for each k' blocks in R do begin**
 - Read k' blocks of R from disk into main memory
 - for each $(k-k')$ block b_s in S do begin**
 - Read $(k-k')$ blocks of S from disk into main memory
 - do a nested loop for each pair (t_r, t_s) [$t_r \in b_r, t_s \in b_s$) to see if they satisfy the join condition *cond*
 - if they do, add $t_r \bullet t_s$ to the result.

Joins – Blocked Nested loop

- Running time depends on amount of main memory buffers available
 - Need at least 2. 1 for R and 1 for S
- Case 4: None of the above (k buffers available between the two tables)
 - Cost: Consider number of blocks read
 - Outer loop: every block in R need to be read once – cost = b_r
 - Inner loop: for each iteration of the outer loop, the whole table in the inner loop need to be read
 - The outer loop executed $\text{ceiling}(b_r / k')$ times
 - So the total number of block reads for inner loop = $b_r * b_s / k'$
 - Given that, what should be the value of k' ?

Joins – Blocked Nested loop

- Running time depends on amount of main memory buffers available
 - Need at least 2. 1 for R and 1 for S
- Case 4: None of the above (k buffers available between the two tables)
 - Cost: Consider number of seeks
 - Outer loop: every time the outerloop executes there need to be a seek (why? – similar to case 2)
 - Number of seeks = $\text{ceiling}(b_r / k')$
 - Inner loop: Similar to case 2
 - Best case scenario = $b_r * (\text{number of seeks for sequential read } S) / k'$
 - Worst case = $b_r * b_s / (k' * (k - k'))$
 - Now what should the value of k' be?

Joins – Blocked Nested loop

- Should we use an index?
- Index for outer loop is useless
 - Unless clustering index on join attribute where you can read the table based on the join attribute, and that attribute is not unique (why?)
- How about inner loop
 - To use an index, that means for each ***tuple*** there need to be a search
 - Number of tuples is larger than number of blocks
 - If each tuples only match with very few tuples, it's fine
 - However, if there are potential large number of matches with a secondary index, things can get dicey.... (exercise)

Joins – Sort-merge

- Consider an equi-join: $R \bowtie_{R.a = S.b} S$
- Assume
 - R is a sequential file ordered by attribute a
 - S is s sequential file ordered by attribute b
- Now to join the two tables
 - We can use the merge algorithm from merge sort

	<i>a1</i>	<i>a2</i>		<i>a1</i>	<i>a3</i>
<i>pr</i> →	a	3	→ <i>ps</i>	a	A
	b	1		b	G
	d	8		c	L
	d	13		d	N
	f	7		m	B
	m	5			
	q	6			
					<i>s</i>
		<i>r</i>			

Joins – Sort-merge

- Key difference
 - If attributes have duplicate values then one may have to go “back and forth”
- Cost for the merge
 - Assume we have k' buffers for R and $k - k'$ buffers for S
 - Blocks read = $b_r + b_s$ (best case, or unique value of attributes)
 - In general, may multiply by some factor m to take account for duplicates
 - Seeks = $\text{ceiling}(b_r / k) + \text{ceiling}(b_s / (k - k'))$ (why?)

Joins – Sort-merge

- What if the tables are not sorted?
- Sort them (!) and store the sorted table temporarily
- Then apply the merge algorithm
- This is known as the sort-merge algorithm
- Cost = cost to sort the tables + merge cost (as previous slide)

Joins – Sort-merge

- Example: Consider joining two tables R (1000 pages), S (100000 pages)
- Assume we have 11 pages of buffers
- Consider sorting R
 - First step: Divide page into segments of 10 pages = 100 segments (to make analysis easier)
 - Merging step: assume merge 10 segment as a time
 - Two iterations: 100 segments, 10 page each -> 10 segments, 100 page each -> sorted, 1000 pages
 - Total number of read/writes = 3 (total iterations) * 2 (read/write) * 1000 = 60000

Joins – Sort-merge

- Example: Consider joining two tables R (1000 pages), S (100000 pages)
- Assume we have 11 pages of buffers
- Consider sorting S
 - First step: Divide page into segments of 10 pages = 10000 segments (similar as R)
 - Merging step: assume merge 10 segment as a time
 - Four iterations : 10000 segments, 10 page each -> 1000 segments, 100 page each -> 100 segments, 1000 page each -> 10 segments, 10000 pages -> sorted, 100000 pages
 - Total number of read/writes = 5 (total iterations) * 2 (read/write) * 100000 = 1000000


Joins – Sort-merge

- Example: Consider joining two tables R (1000 pages), S (100000 pages)
- Assume we have 11 pages of buffers
- Total cost = Cost of sorting R + Cost of sorting S + cost of merge
- $= 60000 + 1000000 + (1000 + 100000) = 1161000$ pages

Joins – Hash joins

- $R \bowtie_{R.a = S.b} S$
- Build a hash table on the file for attribute a and b, using the same hash function for both tables
- Now only tuples in the same buckets in the corresponding table can join together

Joins – Hash joins



Join – Hash joins

- Now each pair of partition can be join by using any algorithm
 - If one of the partition is small enough to fit in the buffers, do a nested loop
 - Otherwise, one can recursively apply hash join – using a different hash function for each recursive call, until one of the partition is small enough

Join – Hash joins

- Cost
 - Harder to estimate
 - Depends a lot on data distribution (as a skewed distribution will lead some segments very long)
 - However, potential advantages
 - As long as one table's tuple behave nicely with the hash table, it should be fine
 - Even if distributions are skewed for both tables, it is possible that they complement each other (i.e. for a hash value, one table may have a lot of tuples, but the other may have little)

Join – Hash joins

- Example: Consider joining two tables R (1000 pages), S (100000 pages)
 - Assume we have 11 pages of buffers
 - Assume all hash function evenly distribute the tuples for both tables
 - First iteration:
 - R and S divided into 10 segments (100 page for R, 10000 page for S)
 - Second iteration
 - Each segment is further subdivided (10 page for R, 1000 page for S)
 - Now the segments for R is small enough to fit in main memory
 - Nested loop for each pair
- Running time = $4 * (100000 + 1000) + 100000 + 1000$
- = 505000
- Better than sort-merge

Joins – hash-join vs. sort-merge

- Comparing sort-merge and hash-join
 - Hash-join is better than sort-merge when there is a large different between number of pages between the tables
 - The number of iteration is dominated by the larger table in sort-merge (need to completely sort both tables)
 - But is dominated by the smaller table in hash-join (need to recursive call until ONE of the segment is small enough)
 - Sort-merge has more predictable performance
 - Hash join's performance depend on how the hash function performs
 - Sort-merge has the output sorted by the join attribute
 - May be important (see later)

Projection

- Seems straightforward – just picking the corresponding values out of the tuple
- But how about SELECT DISTINCT?
 - Need to remove duplicates
- Much trickier than you think
 - Data too large to fit in main memory...

Projection

- Use external sort
 - At each step, eliminate duplicates for a segment before writing on the disk
- Hashing can be used instead

Group by + Aggregation

- **Aggregation** can be implemented in a manner similar to duplicate elimination.
 - **Sorting** or **hashing** can be used to bring tuples in the same group together, and then the aggregate functions can be applied on each group.
 - Optimization: **partial aggregation**
 - combine tuples in the same group during run generation and intermediate merges, by computing partial aggregate values
 - For count, min, max, sum: keep aggregate values on tuples found so far in the group.
 - When combining partial aggregate for count, add up the partial aggregates
 - For avg, keep sum and count, and divide sum by count at the end

Union/Intersect/Except (set operations)

- **Set operations** (\cup , \cap and $-$): can either use variant of sort-merge, or variant of hash-join.
- E.g., Set operations using hashing:
 1. Partition both relations using the same hash function
 2. Process each partition i as follows.
 1. Using a different hashing function, build an in-memory hash index on r_i .
 2. Process s_i as follows
 - $r \cup s$:
 1. Add tuples in s_i to the hash index if they are not already in it.
 2. At end of s_i add the tuples in the hash index to the result.

Union/Intersect/Except (set operations)

- Set operations using hashing:

1. as before partition r and s ,

2. as before, process each partition i as follows

1. build a hash index on r_i

2. Process s_i as follows

- $r \cap s$:

1. output tuples in s_i to the result if they are already there in the hash index

- $r - s$:

1. for each tuple in s_i , if it is there in the hash index, delete it from the index.

2. At end of s_i add remaining tuples in the hash index to the result.

Difference in except: since operation is NOT commutative, restriction on which table can be “outer” loop.

Outer joins

- Challenge: The unmatched tuples still need to be returned (with NULL value padded)
- Sort-merge:
 - Straight-forward: output also unmatched tuples during merge
- Hash-join:
 - Once again, restriction on what can be the “outer” loop