

CS 5330/7330

Two phase locking

Table of contents

- Lock-based protocol
- 2-phase locking
- Deadlocks
- Lock implementation
- The phantom problem

The story so far

- Isolation: one key requirement for transaction
- Isolation as conflict serializability
- Why conflict serializability
 - Cycles in serializability
 - Multiple transactions share the same objects
 - The contention goes in *both* directions

So ...


- Need to provide *protocol* (rules on how data item is accessed) to ensure conflict serializability
- Goal of protocol:
 - To allow access for data items that are not required by multiple transactions
 - For those data items required by multiple transaction, restrict access in some way, or limit it to exclusive access
- Balance between safety and efficiency
 - Too restrictive: little or no concurrency, ineffective
 - Too lenient: leads to inconsistency

Lock-based protocols

- “Exclusive” access \Rightarrow locks
- Each database item is associated with locks
- Transaction must obtain locks before accessing the object
- Transaction must release lock when it finishes.

Lock-based protocols

- Example:



- `Lock(X)`: check if object `X` is already locked
 - If not, obtain the lock
 - If so, wait or “do something” to handle the potential deadlock (like aborting)
- One does not have to read immediately after locking
- `Unlock(X)`: release the lock on object `X`
- The addition of `Lock(X)` and `Unlock(X)` commands are done by the DBMS

Lock-based protocols: S and X locks

- How many transaction can obtain the lock to an item?
 - One
- Too restrictive?
- Consider the two transactions

1. Read(X)
2. Read(Y)
3. Read(Z)

T1

1. Read(Y)
2. Read(X)
3. Read(Z)

T2

- There seems to be no reason for one transaction to wait for the other

Lock-based protocols: S and X locks

- Two kinds of locks on each object
 - Shared locks (S-locks)
 - Requested before reading
 - Multiple transactions can hold a S-lock on an object simultaneously
 - Exclusive locks (X-locks)
 - Requested before writing
 - Only one transaction can hold an X-lock on an object at any given time
 - No other transaction can hold any lock (not even a S-lock) if some transaction has an X-lock

Lock-based protocols: S and X locks

- More on S and X locks
 - A transaction that holds an S-lock on an object can read the object
 - A transaction that holds an X-lock on an object can read and write the objects
- Lock-compatibility table

T1 Request \ T2 holds	S-lock	X-lock
	S-lock	X-lock
S-lock	True	False
X-lock	False	False

Lock-based protocol

- Consider the following examples (again):

```
1. A1 <- Read(X)
2. A1 <- A1 - k
3. Write(X, A1)
4. A2 <- Read(Y)
5. A2 <- A2 + k
6. Write(Y, A2)
```

T1 (Transfer)

```
1. A1 <- Read(X)
2. A1 <- A1 * 1.01
3. Write(X, A1)
4. A2 <- Read(Y)
5. A2 <- A2 * 1.01
6. Write(Y, A2)
```


T2 (Dividend)

Lock-based protocol -- example

- With a lock-based protocol, one possible way T1 is transformed

1. A1 \leftarrow Read(X)
2. A1 \leftarrow A1 - k
3. Write(X, A1)
4. A2 \leftarrow Read(Y)
5. A2 \leftarrow A2 + k
6. Write(Y, A2)

T1 (Transfer)



1. S-lock(X)
2. A1 \leftarrow Read(X)
3. A1 \leftarrow A1 - k
4. X-lock(X)
5. Write(X, A1)
6. X-lock(Y)
7. A2 \leftarrow Read(Y)
8. A2 \leftarrow A2 + k
9. Write(Y, A2)
10. Unlock(X)
11. Unlock(Y)

Lock-based protocol -- example

- Notes from the previous example:
 - Locks must be obtained before read/write can begin
 - If a transaction want to read and write the same object, it can either
 - Obtain an X-lock before reading
 - Obtain an S-lock before reading, then obtain an X-lock before writing (it is not automatically granted)
 - A transaction does not have to release locks immediately after use
 - Good or bad?

Lock-based protocol -- example

- Example of schedule with locks

1. S-lock(X)
2. A1 <- Read(X)
3. Unlock(X)
4. A1 <- A1 - k
5. X-lock(X)
6. Write(X, A1)
7. Unlock(X)
8. S-lock(Y)
9. A2 <- Read(Y)
10. Unlock(Y)
11. A2 <- A2 + k
12. X-lock(Y)
13. Write(Y, A2)
14. Unlock(Y)

T1

1. S-lock(X)
2. A1 <- Read(X)
3. Unlock(X)
4. A1 <- A1 * 1.01
5. X-lock(X)
6. Write(X, A1)
7. Unlock(X)
8. S-lock(Y)
9. A2 <- Read(Y)
10. Unlock(Y)
11. A2 <- A2 * 1.01
12. X-lock(Y)
13. Write(Y, A2)
14. Unlock(Y)

T2

Lock-based protocol -- example

- Example of schedule with locks

1. S-lock(X)
2. A1 \leftarrow Read(X)

3. Unlock(X)

4. A1 \leftarrow A1 - k

5. X-lock(X)

5. X-lock(X)

6. Write(X, A1)

7. Unlock(X)

8.

T1

1. S-lock(X)

2. A1 \leftarrow Read(X)

3. Unlock(X)

4. A1 \leftarrow A1 * 1.01

5. X-lock(X)

5. X-lock(X)

T2

No wait: S-locks

T2 waits

T2 can go ahead

T1 waits

T1 can go ahead

Lock based protocols -- questions

- Does having locks this way guarantee conflict serializability?
- Is there any other requirements in the order/manner of acquiring/releasing locks?
- Does it matter when to acquire locks?
- Does it matter when to release locks?

Lock based protocols – need for a protocol

- Suppose we request a lock immediately before reading/writing
- We release a lock immediately after each read/write
- Does it guarantee serializability?

Lock based protocol – need for a protocol

- But with the following schedule

1. $A1 \leftarrow \text{Read}(X)$
2. $A1 \leftarrow A1 - k$
3. $\text{Write}(X, A1)$

1. $A1 \leftarrow \text{Read}(X)$
2. $A1 \leftarrow A1 * 1.01$
3. $\text{Write}(X, A1)$
4. $A2 \leftarrow \text{Read}(Y)$
5. $A2 \leftarrow A2 * 1.01$
6. $\text{Write}(Y, A2)$

4. $A2 \leftarrow \text{Read}(Y)$
5. $A2 \leftarrow A2 + k$
6. $\text{Write}(Y, A2)$

$X : 100 \rightarrow 50 \rightarrow 50.5; Y : 200 \rightarrow 202 \rightarrow 252; X+Y = 302.5$

Not conflict serializable.

Lock based protocol – need for a protocol

S-lock (X)

1. A1 <- Read(X)

Unlock(X)

2. A1 <- A1 - k

X-lock (X)

3. Write(X, A1)

Unlock(X)

1. A1 <- Read(X)

S-lock (X)

2. A1 <- A1 * 1.01

Unlock(X)

3. Write(X, A1)

Unlock(X), Slock(Y)

4. A2 <- Read(Y)

Unlock(Y)

5. A2 <- A2 * 1.01

X-lock (Y)

6. Write(Y, A2)

Unlock(Y)

S-lock (Y)

4. A2 <- Read(Y)

Unlock(Y)

5. A2 <- A2 + k

X-lock (Y)

6. Write(Y, A2)

Unlock(Y)

Not conflict serializable.

Lock based protocol – need for a protocol

- So the previous technique does not work.
- How about, only releasing a lock after all operations on the object is finished?

Lock based protocol – need for a protocol

S-lock (X)

1. A1 \leftarrow Read(X)

2. A1 \leftarrow A1 - k

X-lock (X)

3. Write(X, A1)

Unlock(X)

1. A1 \leftarrow Read(X)

S-lock (X)

2. A1 \leftarrow A1 * 1.01

3. Write(X, A1)

Unlock(X), Slock(Y)

4. A2 \leftarrow Read(Y)

5. A2 \leftarrow A2 * 1.01

6. Write(Y, A2)

X-lock (Y)

Unlock(Y)

S-lock (Y)

4. A2 \leftarrow Read(Y)

5. A2 \leftarrow A2 + k

X-lock (Y)


6. Write(Y, A2)

Unlock(Y)

Not conflict serializable.

Two-phase locking -- motivation

- What is the problem?
- When a transaction release a lock on an object , that means other transactions can obtain a lock on it.



- In this case, there is contention from T1 to T2
- To ensure serializability, we must ensure there is no conflict from T2 back to T1
- How?

Two-phase locking -- motivation

- Ensure that T1 does not read/write anything that T2 read/write.
 - Unrealistic to check in real life
- What is a sufficient condition then?
- Ensure T1 does not read/write **anything** after releasing the lock!
- \Rightarrow (basic) Two-phase locking

Two phase locking – definition

- The basic two-phase locking (2PL) protocol
 - A transaction T must hold a lock on an item x in the appropriate mode before T accesses x .
 - If a conflicting lock on x is being held by another transaction, T waits.
 - Once T releases a lock, it cannot obtain any other lock subsequently.
- Note: a transaction is divided into two phases:
 - A *growing phase* (obtaining locks)
 - A *shrinking phase* (releasing locks)
- Claim : 2PL ensures conflict serializability

Two phase locking – Notations

- If the operation is immaterial, we will simply use $\text{lock}_i[x]$ and $\text{unlock}_i[x]$ to denote the operations that T_i locks and unlocks x , respectively.
- $o_i[x]$: the execution of the operation o by T_i on x
- $<_s$: the ordering relationship between operations in the schedule S (We omit the subscript s if the context is clear)

Two phase locking – Notations

- 2PL implies
 - $o\text{-lock}_i[x] < o_i[x] < o\text{-unlock}_i[x]$ (if T_i commits)
 - i.e. you obtain a lock, then read/write, then unlock
 - Given two conflicting operations $p_i[x]$ and $q_j[x]$, we have either $p\text{-unlock}_i[x] < q\text{-lock}_j[x]$, or $q\text{-unlock}_j[x] < p\text{-lock}_i[x]$
 - i.e. if there are two conflicting operations, they cannot both hold the lock at the same time
 - For any $p\text{-lock}_i[x]$ and $q\text{-unlock}_i[y]$, $p\text{-lock}_i[x] < q\text{-unlock}_i[y]$
 - i.e. Every transaction is 2-phase (you cannot lock something after you unlock anything)

Two phase locking – Proof


- By contradiction
 - If $T_i \rightarrow T_j$ in $SG(S)$, we have a $p_i[x]$ conflict with a $q_j[x]$ and that $p_i[x] < q_j[x]$ in S
 - $\Rightarrow p\text{-unlock}_i[x] < q\text{-lock}_j[x]$
 - If $T_i \rightarrow T_j \rightarrow T_k$ in $SG(S) \Rightarrow \exists x, y$ such that $\text{unlock}_i[x] < \text{lock}_j[x]$ and $\text{unlock}_j[y] < \text{lock}_k[y]$
 - Since $\text{lock}_j[x] < \text{unlock}_j[y]$, we have $\text{unlock}_i[x] < \text{lock}_k[y]$. By induction, for any path $T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_n$ in $SG(S)$, T_1 performs an unlock before T_n has acquired certain lock.
 - Now, if $SG(S)$ has a cycle: $T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_n \rightarrow T_1$, we have T_1 performs an unlock before T_1 has acquired certain lock. A violation of the 2-phase rule!!

Two phase locking – Serializability

- Lock-point: the point where the transaction obtains all the locks
- With 2PL, a schedule is conflict equivalent to a serial schedule ordered by the lock-point of the transactions

2-phase locking -- example

1. S-lock(X)
2. $A1 \leftarrow \text{Read}(X)$
3. $A1 \leftarrow A1 - k$
4. X-lock(X)
5. $\text{Write}(X, A1)$
6. S-lock(Y)
7. $A2 \leftarrow \text{Read}(Y)$
8. $A2 \leftarrow A2 + k$
9. X-lock(Y)
10. $\text{Write}(Y, A2)$
11. Unlock(X)



12. Unlock(Y)

T1

1. S-lock(X)



T2 waits

1. S-lock(X)
2. $A1 \leftarrow \text{Read}(X)$
3. $A1 \leftarrow A1 * 1.01$
4. X-lock(X)
5. $\text{Write}(X, A1)$
6. S-lock(Y)



T2 waits

6. S-lock(Y)
7. $A2 \leftarrow \text{Read}(Y)$
8. $A2 \leftarrow A2 * 1.01$
9. X-lock(Y)
10. $\text{Write}(Y, A2)$
11. Unlock(Y)
12. Unlock(X)



Lock point for T2

T2

2 phase locking -- recoverability

- Recall definitions for recoverability
 - Recoverability
 - Avoid cascade aborts
 - Strict
- If a schedule is 2PL, does it guarantee any of the above?

2 phase locking -- recoverability

- Consider the following schedule

1. X-lock(X)
2. A1 <- Read(X)
3. A1 <- A1 * 10
4. Write(X, A1)
5. Unlock(X)

1. X-lock(X)
2. A2 <- Read(X)
3. A2 <- A2 + 1
4. Write(X, A2)
5. Unlock(X)
6. Commit

6. Abort!

Not recoverable!

2 phase locking – recoverability

- What's the problem?
- There is a gap between releasing locks and the decision to commit/abort
- Other transactions can still access data written by a uncommitted transaction
- How to solve this problem?

2 phase locking – recoverability

- Strict 2-phase locking:
 - 2-phase locking
 - X-locks can only be released when the transaction commits
- Question: does strict 2-phase locking ensure
 - Recoverability?
 - Avoid cascade aborts?
 - Strict?

2 phase locking – recoverability

- Rigorous 2-phase locking:
 - 2-phase locking
 - Any lock can only be released when the transaction commits
- Ensure serializability
 - Moreover, serial schedule is ordered by the time the transaction commits

Lock conversion (upgrades)

- Consider the transaction on the right
- Eventually transaction will need an X-lock on X
- But obtaining it at the beginning seems a waste of effort
- \Rightarrow lock conversion
- Obtain an S-lock on X first, then obtain an X-lock only at the end
- Notice that X-lock is not automatically granted
- Similarly, downgrades are allowed (X-lock \rightarrow S-lock)

1. Read(X)
2. Read(Y)
3. Read(Z)
4. ...

1000 operations not involving X
1005. Write(X)

S-lock(X)

X-lock(X)

Lock conversion (upgrades)

- Why lock conversion?
 - Allow more concurrency
 - If the lock operations are generated automatically by CC, we have no information about whether a transaction that reads x will or will not write x later.
- Upgrades are consider equivalent as obtaining a new lock (must be in growing phase)
- Downgrades are consider equivalent as releasing a lock (must be in shrinking phase)

Deadlocks

- 2 phase locking is a blocking protocol (transaction has to wait if it cannot obtain a lock)
- Probability of a deadlock
- Example:

1. S-lock(X)
2. $A1 \leftarrow \text{Read}(X)$
3. $A1 \leftarrow A1 - k$

Can't proceed T2
has S-lock on X

4. X-lock(X)

1. S-lock(X)
2. $A1 \leftarrow \text{Read}(X)$
3. $A1 \leftarrow A1 * 1.01$

Can't proceed T1
has S-lock on X

4. X-lock(X)

Deadlock!

Deadlocks

- How to deal with deadlocks?
 - Ostrich
 - Pretend nothing happens and wait for user to hit the <ctrl-alt-del> key
 - Timeout
 - Assume deadlock after there is no progress for some time, and do something about it
 - Detection and recovery
 - Wait until a deadlock occurs and do something about it
 - Avoidance
 - Wait until a deadlock can occur if certain operations is executed and do something about it
 - Prevention
 - Set up the system such that there is never a chance of deadlocks

Deadlocks -- timeout

- When transaction waits too long for a lock, it is aborted and restarted again from the beginning
- The timeout period should be:
 - long enough so that most transactions that are aborted are actually deadlocked;
 - short enough so that deadlocked transactions don't wait too long for their deadlocks to be broken.

Deadlocks -- Detection

- Deadlocks can be described as a *wait-for graph*, which consists of a pair $G = (V, E)$,
 - V is a set of vertices (all the transactions in the system)
 - E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$.
- If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j , implying that T_i is waiting for T_j to release a data item.
- When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted in the wait-for graph. This edge is removed only when T_j is no longer holding a data item needed by T_i .
- The system is in a deadlock state if and only if the wait-for graph has a cycle. Must invoke a deadlock-detection algorithm periodically to look for cycles.

Deadlocks -- Detection

- When deadlock is detected :
 - Some transaction will have to rolled back (made a victim) to break deadlock. Select that transaction as victim that will incur minimum cost.
 - Rollback -- determine how far to roll back transaction
 - **Total rollback**: Abort the transaction and then restart it.
 - More effective to roll back transaction only as far as necessary to break deadlock. (but tricky to implement)
- Drawback:
 - maintaining wait-for-graph is not cheap
 - Detecting cycles is an overhead

Deadlocks -- Starvation

- Sometimes the same transaction may keep being aborted
- This leads to starvation (really no blocking, but one transaction never get progressed)
- Solution: incorporate number of times being rolled-back as part of the cost

Deadlocks -- Avoidance

- Allow normal operation, but when there is a chance that deadlock will occur, then do something about it
- Recall that deadlock cannot occur if the wait-for-graph have no cycles.
- If we numbered each transaction (e.g. timestamp), then
 - If every edge in the wait-for graph points from $T_i \rightarrow T_j$ such that $i < j$, then deadlock never occurs
 - Why?
 - Works also if for every edge, $i > j$

Deadlocks -- Avoidance

- When a transaction T_i request a lock on an object, but T_j currently have the lock
- Two options:
 - Wait-die (non pre-emptive)
 - If $i < j$ then wait
 - else T_i aborts
 - Wound-wait (pre-emptive)
 - If $i > j$ then wait
 - Else T_j aborts
- If i and j represent time (i.e. small number = older transactions)
 - Wait-die: older transactions wait for younger transactions;
 - Wound-wait: younger transaction wait for older transactions

Deadlocks -- Avoidance

- How to avoid starvation?
- Note that in both scheme, at any given time, at least one transaction is never going to be aborted
 - The one with the smallest i
- Thus, to avoid starvation, ensure that when a transaction is aborted, it is restarted with the SAME i
 - Eventually it will becomes the one with the smallest i

Deadlock -- prevention

- Select a scheme such that deadlock will never occurs
- Conservative 2PL
 - 2 phase locking
 - A transaction must request ALL the locks in the beginning, and either all the locks or none of the locks are granted
- Labelling database objects
 - Transaction can only request locks in order of the database objects

Locks -- implementation

- Various support need to implement locking
 - OS support – lock(X) must be an atomic operation in the OS level
 - i.e. support for critical sections
 - Implementation of read(X)/write(X) – automatically add code for locking
 - Lock manager – module to handle and keep track of locks

Locks -- implementation

- A transaction T_i issues the standard read/write instruction, without explicit locking calls.
- The operation **read**(D) is processed as:

```
    if  $T_i$  has a lock on  $D$ 
    then
        read( $D$ )
    else
        begin
            if necessary wait until no other
                transaction has a X-lock on  $D$ 
            apply any deadlock avoidance rules)
            grant  $T_i$  a S-lock on  $D$ ;
            read( $D$ )
        end
```

(or

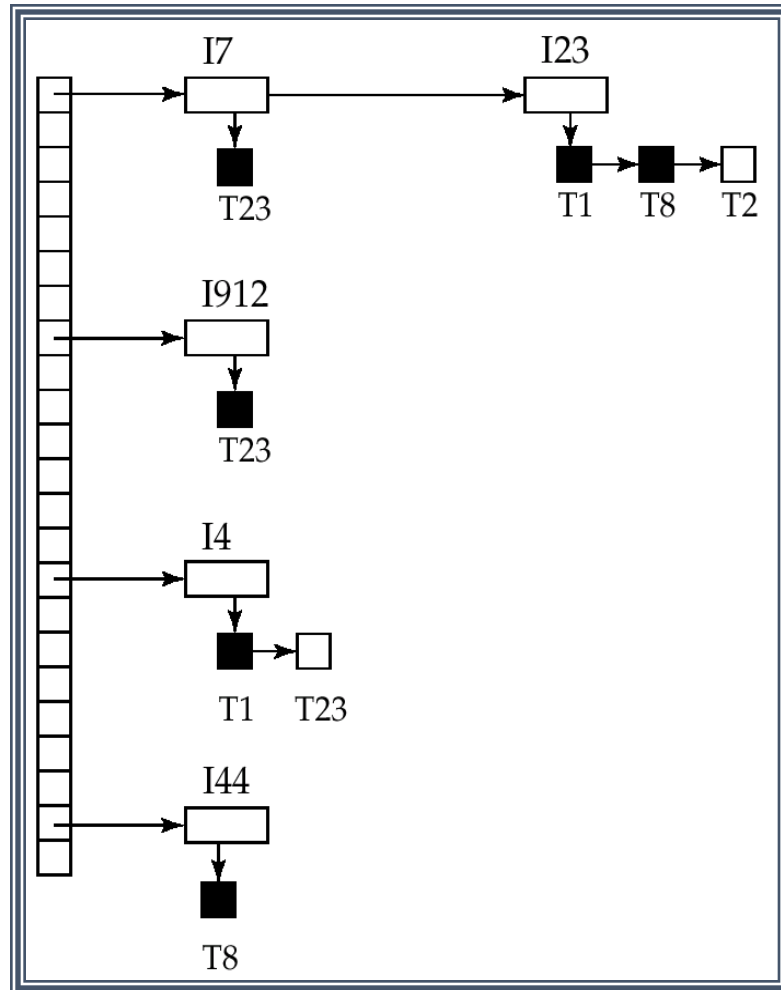
Locks -- implementation

- **write(D)** is processed as:
 - if T_i has a **X-lock** on D
 - then**
 - write(D)
 - else**
 - begin**
 - if necessary wait until no other trans. has any lock on D (or apply any deadlock avoidance rules),
 - if T_i has a **S-lock** on D
 - then**
 - upgrade** lock on D to **X-lock**
 - else**
 - grant T_i a **X-lock** on D
 - write(D)
 - end;**
 - All locks are released after commit or abort

Locks – Lock manager

- A **Lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The lock manager maintains a data structure called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked

Locks – lock manager



- Black rectangles indicate granted locks, white ones indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
 - lock manager may keep a list of locks held by each transaction, to implement this efficiently

Insertion & Deletion

- Does delete(X) conflict with read(X)/write(X)?
 - Yes. Wrong order leads to errors
- Thus delete(X) should be treated like a write operation
 - Request X-locks
- Similar to insert(X) operation.
 - X-lock is given to the newly created tuple.

The phantom menace

- Consider the following 2 transactions

1. Select sum(salary)

From faculty

Where dept = "CS"

2. Select sum(salary)

From faculty

Where dept = "Math"

T1

a. Insert into faculty

values ("Lin", "CS", 1000)

b. Insert into faculty

values ("Lam", "Math", 5000)

T2

- There does not seem to be a conflict (in terms of tuple)
- Assume initially CS faculty have total salary 1,000,000 and Math faculty have total salary 2,000,000
- Then T1 -> T2 will imply the select statements return 1,000,000 and 2,000,000
- T2 -> T1 will imply the select statements return 1,001,000 and 2,005,000

The phantom menace

- But consider the following schedule

1. Select sum(salary)
From faculty
Where dept = "CS"
2. Select sum(salary)
From faculty
Where dept = "Math"

T1

a. Insert into faculty
values ("Lin", "CS", 1000)

b. Insert into faculty
values ("Lam", "Math", 5000)

T2

- The output will be 1,000,100 and 2,000,000
- Not conflict serializable!

The phantom problem

- This is known as *the phantom problem*
- Why does it occur?
 - No tuples are in conflict
 - However, conflict occurs for tuples that satisfy a certain condition (dept = “CS”, dept = “Math”)
 - T1 require access for ALL tuples satisfying the condition
 - However, T2 changes the number of tuples satisfying the condition
- No quick solution: index-locking as a possibility (next lecture)